

Group Meeting

Introduction to numerical optimization and derivatives

Doug Shi-Dong

Computational Aerodynamics Group
Department of Mechanical Engineering
McGill University

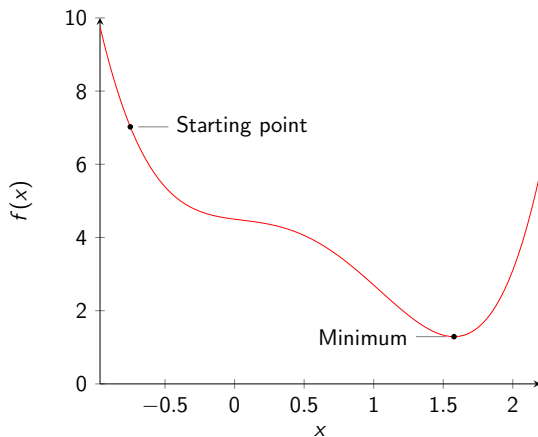


24 March 2021

Problem statement

- Minimize an objective function $f : \mathbb{R}^n \rightarrow \mathbb{R}$



$$\min_x f(x)$$



Minimum

- What defines a minimum?
- First-order necessary condition

$$\frac{df}{dx} = 0 \quad (1)$$




- Second-order sufficient condition

$$v^T \frac{d^2f}{dx^2} v > 0 \quad (2)$$

Direction of descent

- We don't know the objective landscape
- We only know "local" information



Gradient descent

- Evaluate gradient the current design and step in its negative direction

$$x^{n+1} = x^n - \frac{df}{dx} \quad (3)$$

- Gradient is the local slope, therefore, must step carefully

$$x^{n+1} = x^n - \eta \frac{df}{dx} \quad (4)$$

where η is also known as the step length, or learning rate (for ML).

- However, as we approach our minimum $\frac{df}{dx} \rightarrow 0$, which means that we take smaller and smaller steps


Newton step

- Model f as a quadratic function


$$f \approx (\Delta x)^T \frac{d^2 f}{dx^2} (\Delta x) + \frac{df}{dx} (\Delta x) + f \quad (5)$$

- Easy to find minimum of a quadratic function, "simply" solve for (Δx)


$$\frac{d^2 f}{dx^2} (\Delta x) = -\frac{df}{dx} \quad (6)$$




Newton step




Newton step



Newton step



Newton step



- Compute derivatives various ways
 - Analytical
 - Finite differences
 - Automatic differentiation

	Analytical	Finite-Differences	Automatic Differentiation
Easy to implement	NO	YES	YES
Sustainable code	NO	YES	YES
Exact	YES	NO	YES
Scalable	YES	NO	YES

Analytical derivatives

- Good for small functions or to assemble blocks of derivatives
- Tedious and error-prone

	Analytical
Easy to implement	NO
Sustainable code	NO
Exact	YES
Scalable	YES

Analytical derivatives

[illegible]

Figure: 100 out of 2800 lines of boundary conditions for quasi-1D Euler

Finite differences derivatives

- Need to choose a good perturbation
- Need as many function evaluations as design variables




Figure: Finite difference error

	Analytical	Finite-Differences
Easy to implement	NO	YES
Sustainable code	NO	YES
Exact	YES	NO
Scalable	YES	NO

Automatic differentiation derivatives

- It's automatic right?

	Analytical	Finite-Differences	Automatic Differentiation
Easy to implement	NO	YES	YES
Sustainable code	NO	YES	YES
Exact	YES	NO	YES
Scalable	YES	NO	YES

Automatic differentiation derivatives

- It's automatic right? Yes. But no.
- Needs code planification and a bit of software engineering

	Analytical	Finite-Differences	Automatic Differentiation
Easy to implement	NO	YES	YES (but no)
Sustainable code	NO	YES	YES
Exact	YES	NO	YES
Scalable	YES	NO	YES (if done correctly)

C++ templates

```
double function(double input)
{
    double result = std::sin(input) * input;
    return result;
}

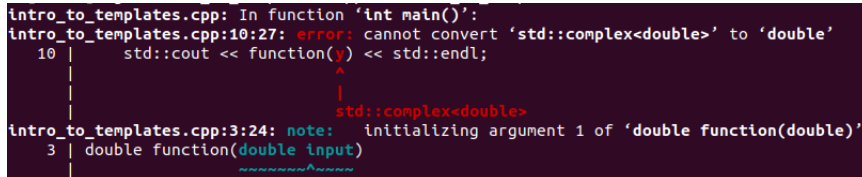
int main() {
    double x = 3.0;
    std::cout << function(x) << std::endl;
}
```

C++ templates

```
double function(double input)
{
    double result = std::sin(input) * input;
    return result;
}

int main() {
    double x = 3.0;
    std::cout << function(x) << std::endl;

    std::complex<double> y(3.0, 1.0);
    std::cout << function(y) << std::endl;
}
```



```
intro_to_templates.cpp: In function 'int main()':
intro_to_templates.cpp:10:27: error: cannot convert 'std::complex<double>' to 'double'
  10 |         std::cout << function(y) << std::endl;
      |                             ^
      |                             |
      |                         std::complex<double>
intro_to_templates.cpp:3:24: note: initializing argument 1 of 'double function(double)'
   3 | double function(double input)
      |
```

Figure: Resulting compilation error

C++ templates

```
template<typename realtype>
realtype function(realtype input)
{
    realtype result = std::sin(input) * input;
    return result;
}

int main() {
    double x = 3.0;
    std::cout << function<double>(x) << std::endl;

    std::complex<double> y(3.0, 1.0);
    std::cout << function<std::complex<double>>(y) << std::endl;
}
```

C++ templates used for operator overloading

```
template<typename realtype>
double function(double input)
{
    double result = std::sin(input) * input;
    return result;
}

std::complex<double> function(std::complex<double> input)
{
    std::complex<double> result = input*input;
    return result;
}

int main() {
    double x = 3.0;
    std::cout << function(x) << std::endl;

    std::complex<double> y(3.0, 1.0);
    std::cout << function(y) << std::endl;
}
```

Automatic differentiation

- Automatic differentiation define their own type such as
 - Sacado::Fad::DFad<double> for Sacado
 - codi::RealForward for CoDiPack.

```
template<typename realtype>
realtype function(realtype input)
{
    realtype result = std::sin(input) * input;
    return result;
}

int main() {
    double x = 3.0;
    std::cout << function<double>(x) << std::endl;

    codi::RealForward y = x; // y = (3.0, 0.0)
    y.setGradient(1.0); // y = (3.0, 1.0)
    // Automatic diff, chain rule
    // df/dinput = dsin(input)/dinput * dinput/dinput * input
    //              + sin(input) * dinput/dinput
    // f = (sin(3.0)*3.0, cos(input)*1.0*3.0 + sin(3.0)*1.0) = (0.42336, -2.8289)
    codi::RealForward f = function<codi::RealForward>(y);
    double gradient = f.getGradient();
    std::cout << f << std::endl; // Outputs 0.4234
    std::cout << gradient << std::endl; // Outputs -2.8289
}
```

Automatic differentiation

- In the forward-mode, imagine that the input variable of type `Sacado::Fad::DFad<double>` is represented by a set of two doubles. $x = (3.0, 1.0)$.
- When a function is called on it, for example $z = \sin(x)$, the AD libraries know that it should be storing $z = (\sin(3.0), 1.0 * \cos(3.0))$
- Once all operators, such as $(*, /, +, -, \text{pow}, \text{sqrt}, \text{etc.})$ have their differentiator counterpart defined by the AD library, it is easy to imagine how to chain those operations.

Automatic differentiation

- Nocedal 2006

$$f(x) = (x_1 x_2 \sin x_3 + e^{x_1 x_2}) / x_3 \quad (7)$$




Figure 8.2 Computational graph for $f(x)$ defined in (8.26).

Forward-mode and Reverse-mode

- Forward-mode cost is proportional to the number of **inputs**
- Reverse-mode cost is proportional to the number of **outputs**
- Reverse-mode requires a "tape" that stores operations and runs it "backwards".
- Gradient of a single objective function is a great use of reverse-mode

- Easier to show
- Implementation will depend on AD library being used
- <https://github.com/dougshidong/OptimizationTutorial>